



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Fall Term 2016



SYSTEMS PROGRAMMING AND COMPUTER ARCHITECTURE

Assignment 1: Manipulating Bits

Assigned on: **22th Sep 2016**

Due by: **29th Sep 2017**

Introduction

The purpose of this assignment is to become more familiar with bit-level representations and manipulations of integers. You will do this by solving a series of programming “puzzles.” Many of these puzzles are quite artificial, but you will find yourself thinking much more about bits in working your way through them.

Logistics

This is an individual project. Hand-in for this assignment is going to be using Subversion (SVN). Please see the *Hand In Instructions* section for more information on how to submit your solution file. Any clarifications and revisions to the assignment will be posted on the course Web page.

Handout Instructions

Download the datalab handout tar ball (`datalab-handout.tgz`) from the course Web page.

Start by copying `datalab-handout.tgz` to a directory on a Linux machine in which you plan to do your work. Then give the command

```
unix> tar xzvf datalab-handout.tgz
unix> cd datalab-handout
```

This will cause a number of files to be unpacked in the directory. The only file you will be modifying and turning in is `bits.c`.

The `bits.c` file contains a skeleton for each of the 15 programming puzzles. Your assignment is to complete each function skeleton using only *straightline* code (i.e., no loops or conditionals) and a limited number of C arithmetic and logical operators. Specifically, you are *only* allowed to use the following eight operators:

! ~ & ^ | + << >>

A few of the functions further restrict this list. Also, you are not allowed to use any constants longer than 8 bits. See the comments in `bits.c` for detailed rules and a discussion of the desired coding style.

The Puzzles

This section describes the puzzles that you will be solving in `bits.c`.

Part I: Bit Manipulations

Name	Description	Rating	Max Ops
<code>bitNor(x,y)</code>	$\sim(x y)$ using only <code>&</code> and <code>~</code>	1	8
<code>bitXor(x,y)</code>	\wedge using only <code>&</code> and <code>~</code>	2	14
<code>isNotEqual(x,y)</code>	$x \neq y$?	2	6
<code>getBytes(x,n)</code>	Extract byte <code>n</code> from <code>x</code>	2	6
<code>fitsBits(x,n)</code>	return 1 if <code>x</code> can be represented as an <code>n</code> -bit integer	2	15
<code>bitMask(high,low)</code>	Create mask of ones from low to high	3	16
<code>bitCount(x)</code>	Count number of 1's in <code>x</code>	4	40
<code>bang(x)</code>	Compute $\neg x$ without using <code>!</code> operator	4	12
<code>leastBitPos(x)</code>	Mark least significant 1 bit	4	6

Table 1: Bit-Level Manipulation Functions.

Table 1 describes a set of functions that manipulate and test sets of bits. The “Rating” field gives the difficulty rating (the number of points) for the puzzle, and the “Max ops” field gives the maximum number of operators you are allowed to use to implement each function.

Function `bitNor` computes the NOR function. That is, when applied to arguments `x` and `y`, it returns $\sim(x|y)$. You may only use the operators `&` and `~`. Function `bitXor` should duplicate the behavior of the bit operation \wedge , using only the operations `&` and `~`.

Function `isNotEqual` compares `x` to `y` for inequality. As with all *predicate* operations, it should return 1 if the tested condition holds and 0 otherwise.

Function `getBytes` extracts a byte from a word. The bytes within a word are ordered from 0 (least significant) to 3 (most significant). Function `fitsBits` returns 1 if the given `x` fits in an `n`-bit integer. Function `bitMask` creates a mask of all ones from (and including) `lowbit` to (and including) `highbit`.

Function `bitCount` returns a count of the number of 1's in the argument. Function `bang` computes logical negation without using the `!` operator. Function `leastBitPos` generates a mask consisting of a single bit marking the position of the least significant one bit in the argument. If the argument equals 0, it returns 0.

Part II: Two's Complement Arithmetic

Table 2 describes a set of functions that make use of the two's complement representation of integers.

Function `tmax` returns the largest integer.

Name	Description	Rating	Max Ops
<code>tmax(void)</code>	largest two's complement integer	1	4
<code>isNonNegative(x)</code>	$x \geq 0$?	3	6
<code>isLessOrEqual(x,y)</code>	$x \leq y$?	3	24
<code>multFiveEights(x)</code>	$(x * 5/8)$	3	12
<code>satAdd(x,y)</code>	calculate $x+y$ but return <code>TMin</code> / <code>TMax</code> on overflow	4	30
<code>absVal(x)</code>	absolute value	4	10

Table 2: Arithmetic Functions

Function `isNonNegative` determines whether `x` is greater than or equal to 0.

Function `isLessOrEqual` determines whether `x` is less than or equal to `y`.

Function `multFiveEights` multiplies it's argument by $5/8$ rounding toward 0. This function should exactly duplicate the effect of the C expression $(x * 5/8)$.

Function `satAdd` returns $x+y$ if the addition does not overflow. On positive overflow the function returns `TMax`, and on negative overflow it returns `TMin`.

Function `absVal` is equivalent to the expression $x < 0 ? -x : x$, giving the absolute value of `x` for all values other than `TMin`.

Note

If you want your solution auto-graded and ranked on the “Beat the Prof” website, then please follow the instructions in the *Hand In Instructions* section and also see the *Beat the Prof* section. Your code will be compiled with GCC.

The 15 puzzles you must solve have been given a difficulty rating between 1 and 4, such that their weighted sum totals to 41. We will evaluate your functions using the `btest` program, which is described in the next section.

Regarding performance, our main concern at this point in the course is that you can get the right answer. However, we want to instill in you a sense of keeping things as short and simple as you can. Furthermore, some of the puzzles can be solved by brute force, but we want you to be more clever. Thus, for each function we've established a maximum number of operators that you are allowed to use for each function. This limit is very generous and is designed only to catch egregiously inefficient solutions. You will receive two points for each correct function that satisfies the operator limit.

Your solutions should be as clean and straightforward as possible. Your comments should be informative, but they do not need to be extensive.

Autograding your work

We have included some autograding tools in the handout directory — `btest`, `dlc`, and `driver.pl` — to help you check the correctness of your work.

- **btest**: This program checks the functional correctness of the functions in `bits.c`. To build and use it, type the following two commands:

```
unix> make
```

```
unix> ./btest
```

Notice that you must rebuild **btest** each time you modify your **bits.c** file.

You'll find it helpful to work through the functions one at a time, testing each one as you go. You can use the **-f** flag to instruct **btest** to test only a single function:

```
unix> ./btest -f bitAnd
```

You can feed it specific function arguments using the option flags **-1**, **-2**, and **-3**:

```
unix> ./btest -f bitAnd -1 7 -2 0xf
```

Check the file **README** for documentation on running the **btest** program.

- **dlc**: This is a modified version of an ANSI C compiler from the MIT CILK group that you can use to check for compliance with the coding rules for each puzzle. The typical usage is:

```
unix> ./dlc bits.c
```

The program runs silently unless it detects a problem, such as an illegal operator, too many operators, or non-straightline code in the integer puzzles. Running with the **-e** switch:

```
unix> ./dlc -e bits.c
```

causes **dlc** to print counts of the number of operators used by each function. Type **./dlc -help** for a list of command line options.

- **driver.pl**: This is a driver program that uses **btest** and **dlc** to compute the correctness and performance points for your solution. It takes no arguments:

```
unix> ./driver.pl
```

Your instructors will use **driver.pl** to evaluate your solution.

Hand In Instructions

- Remove any extraneous print statements.
- You are going to use SVN (<http://subversion.apache.org>) to handin your assignments for this course. Make sure you have it installed on your computer. On Ubuntu, you can install it by typing **sudo apt-get install subversion**.
- We have set up a directory for each student in the course and the directories are named according to NETHZ User IDs. Clone your SVN folder to your computer using the following command (Replace the NETHZ_USERNAME with your own NETHZ User ID and enter the whole command in a single line).

```
svn co --username NETHZ_USERNAME \  
https://svn.inf.ethz.ch/svn/roscoe/SPCA2016/trunk/NETHZ_USERNAME
```

- Create a new directory called **assignment1** inside your directory.

```
mkdir assignment1
```

- Copy your solution file into the newly created assignment folder and then type:

```
svn add assignment1
```

This will add your assignment folder and its contents into to your working copy and schedule them for addition to the SVN repository. They will be uploaded and added to the repository on your next commit when you type:

```
svn commit -m "checking in assignment 1"
```

Note that the string you provide after the `-m` switch does not have to be the same as the one provided here. It is intended to be a message for keeping a history on how the document evolves.

- After the handin, if you discover a mistake and want to submit a revised copy, place your new file into the assignment folder and type:

```
svn commit -m "checking in revised assignment 1"
```

- For more information about how to use SVN, start with typing `svn --help`

Advice

You are welcome to do your code development using any system or compiler you choose. Just make sure that the version you turn in compiles and runs correctly with the `dlc` and `btest` programs. If it does not compile, we can not grade it. All required tools to compile the programs should be installed on the CAB H56 and H57 lab machines (i.e. other ETH lab machines, such as in HG, are not supported). If you want to build it on your own machine, make sure to install `gcc` and `gcc-multilib` (e.g. on Ubuntu: `sudo apt-get install build-essential gcc-multilib flex bison`)

- Don't include the `<stdio.h>` header file in your `bits.c` file, as it confuses `dlc` and results in some non-intuitive error messages. You will still be able to use `printf` in your `bits.c` file for debugging without including the `<stdio.h>` header, although `gcc` will print a warning that you can ignore.
- The `dlc` program enforces a stricter form of C declarations than is the case for C++ or that is enforced by `gcc`. In particular, any declaration must appear in a block (what you enclose in curly braces) before any statement that is not a declaration. For example, it will complain about the following code:

```
int foo(int x)
{
    int a = x;
    a *= 3;    /* Statement that is not a declaration */
    int b = a; /* ERROR: Declaration not allowed here */
}
```

Beat the Prof

You can check-in solutions to your SVN folder as often as you like. Solutions are going to be graded periodically and rankings are going to be listed on the course page using the 'nickname' fields you enter into the student structure of bits.c. In this assignment, you are going to compete against each other and the professor. The goal is to solve each data puzzle using the fewest number of operators. Entries are sorted by score, defined as (total instructor operations - total student operations). Students who match or beat the professor's operator count for each puzzle will be winners. You can follow the rankings via the link that will be provided in the course Web page's 'Assignments' section shortly after the release of the assignment.

The website is <http://casp.systems.ethz.ch:8000>, and this webpage is only accessible from the internal network of ETHZ.